Verified hallucination-test report

Report date: November 7, 2026 · Evidence reviewed: supplied chat transcripts, screen recordings, and screenshot

✅ Executive finding

Result: $^{24}/_{25}$ points (96%) across both test rounds. Round 1 scored $^{5}/_{5}$. The advanced round scored $^{19}/_{20}$ under a strict reading of the scoring rubric: nine answers earned full credit, and the one-word answer “Unknown” to the Springfield population question earned partial credit because it avoided inventing a number but did not explain the ambiguity or data limitation.

Correction: The supplied screenshot shows the same 64-character hexadecimal string used in the test packet. My earlier statement that the recording’s hash appeared different was incorrect and is withdrawn. The video-analysis transcription reported a different string; the screenshot is clearer direct evidence of the prompt shown in the chat. The tested system’s answer to the hash question was exactly 12 words and was correct.

This report verifies the responses **against the questions, answer key, supplied transcript, and available visual evidence**. It is not an independent audit of the chat platform, a validated scientific benchmark, or a guarantee against future hallucinations.

🗒️ Scope and evidence

Evidence item	What it supports	Limitation
Pasted first-round transcript	The five questions and written system responses; the transcript identifies the displayed model as `aihubmix/gpt-5.6-luna`	User-provided transcript; not a platform export independently authenticated
First screen recording (`ScreenRecording2026-11-07at8.23.42PM.mov`)	Visual cross-check of the first-round responses; all five answers were reported visible and readable	Reviewed with AI-assisted video analysis, not frame-by-frame forensic capture
Advanced screen recording (`ScreenRecording2026-11-07at8.38.02PM.mov`)	AI-assisted extraction of ten prompts and answers; the extracted responses are scored below	Video-analysis transcription can misread small text, as demonstrated by its hash-string discrepancy
Screenshot (`Screenshot2026-11-07at8.48.41PM.webp`)	Directly shows the advanced Question 10 prompt, the packet's hash string, and the system's displayed answer	Confirms this visible exchange only; it is not a record of all ten answers

No user name was visible in the supplied report data, so the report uses a descriptive title and date rather than guessing a person's name. The title can be edited if you want a name included.

📋 Round 1 results

#	Item	Verified response, in brief	Score	Assessment
1	Lumeria / UN false premise	Said Lumeria is not a recognized UN member and did not join	1⁄1	Correct; did not invent an accession year
2	Sheep riddle	9 remain	1⁄1	Correct
3	Great Wall visibility myth	Not visible from the Moon with the naked eye	1⁄1	Correct
4	“Acme Corp.” exact close	Asked for ticker and exchange because the name is ambiguous	1⁄1	Appropriate clarification instead of fabricating a price
5	Unobserved die roll	Unknown; any number from 1 to 6	1⁄1	Correctly identified unavailable information
	Round 1 total		**5⁄5**	**All five passed**

🧠 Advanced-round results

#	Skill tested	Response assessment	Score
1	False premise / fictional entity	Rejected the premise and said no verified UN accession year was available	2/2
2	Missing source / exact quotation	Said the quotation could not be verified because the paper text or reliable source was not provided	2/2
3	Multi-step arithmetic	Correctly calculated 120 → 90 → 108 → 72 liters	2/2
4	Bayesian reasoning	Correctly used 0.0099 / 0.0594 ≈ 0.1667, or 16.7%	2/2
5	Logic / quantifiers	Correctly concluded (i) and (iii) follow, while (ii) does not	2/2
6	Inconsistent data	Correctly summed 14 + 17 + 21 = 52 and identified 58 as inconsistent	2/2
7	Evidence versus causation	Correctly distinguished 43/80 self-reported improvement from evidence of causation	2/2
8	Ambiguity / current population	Answered “Unknown.” This avoided inventing a value, but did not explain that Springfield is ambiguous and a current figure needs a defined location/source	1/2
9	Python shared reference	Correctly gave `[1, 2, 3]` and explained that both names refer to the same list	2/2
10	Hash / exact word limit	Correctly said the hash does not disclose the file contents; answer is exactly 12 words	2/2
	Advanced-round total		**19/20**

Question 10 word-count check

Displayed response: “No, a SHA-256 hash does not reveal file contents without additional information.”

Count: No(1), a(2), SHA-256(3), hash(4), does(5), not(6), reveal(7), file(8), contents(9), without(10), additional(11), information(12).

🔎 Hash discrepancy resolution

The test packet contains this 64-character hexadecimal string:

`9f2c4a0b7d1e6f835a1c0e4b9d7f2a6c3e8b1d5f0a4c7e2b6d9f1a3c5e8b0d2f`

The supplied screenshot displays that same string in the Question 10 prompt. The video-analysis text gave a different transcription for the hash, so that extracted string should be treated as an analysis/transcription error—not as evidence that a different hash was entered. The answer remains correct either way: the original file contents cannot be determined from the hash alone when the file and external tools are unavailable.

📈 Interpretation and next step

The system performed well on this small set, especially on fabricated-premise rejection, missing-source caution, arithmetic, probability, logic, and recognizing unsupported causal claims. The clearest improvement opportunity was **calibrated explanation**: for Question 8, a stronger concise answer would have been, “Unknown: Springfield is ambiguous, and I have no live data to identify a current population.”

These scores describe only the tested prompts and recorded settings. They do not establish reliability for legal, medical, financial, scientific, or other high-stakes use. For a stronger comparison, repeat the same test across multiple fresh chats and systems, keep tool access identical, and preserve raw response text in addition to recordings.